

# WiscAFS

Asket Agarwal

Preetham Srinivas Kikkeri

Aanandita Dhawan

## DESIGN

- The file system is implemented with an AFS-like protocol.
- Upon file open, the client will request the entire file from the server and store it in the client's local cache.
- Subsequent read and write requests will be redirected to the local cached copy. On close, the dirty file will be flushed back to the server. If the file was opened in read-only mode then it is not pushed to server.
- Client replicates the server directory structure.
- Server is stateless in our design.

### Client Design (cache)

- Client does a whole file caching where it fetches the files from the server.
- Whenever a user opens a file, the client does a getattr call to the server to see if the file was modified. If yes, it fetches it.
- All reads/writes happens to the local copy with no server interaction to increase performance.
- Upon close the entire file is flushed to the server.

## Filebench Metrics

### Filecreate.f

```
finish          432ops      43ops/s   0.0mb/s    0.000ms/op [0.000ms - 0.001ms]
append-file     433ops      43ops/s   43.2mb/s   23.036ms/op [3.648ms - 25.253ms]
11.050: IO Summary:  433 ops 43.296 ops/s 0/43 rd/wr  43.2mb/s 23.036ms/op
```

### Filemicro\_createfiles

```
finish          2436ops     244ops/s   0.0mb/s    0.000ms/op [0.000ms - 0.001ms]
closefile1      2436ops     244ops/s   0.0mb/s    0.134ms/op [0.086ms - 0.257ms]
writefile1      2436ops     244ops/s   0.2mb/s    0.333ms/op [0.186ms - 0.511ms]
createfile1     2436ops     244ops/s   0.0mb/s    3.608ms/op [2.964ms - 6.718ms]
409.681: IO Summary: 7308 ops 730.716 ops/s 0/244 rd/wr  0.2mb/s 1.358ms/op
```

### Filemicro\_createrand

```
11.248: Per-Operation Breakdown
finish          80ops       8ops/s    0.0mb/s    0.000ms/op [0.000ms - 0.002ms]
sync            80ops       8ops/s    0.0mb/s   15.627ms/op [8.553ms - 132.780ms]
append-file     808ops      81ops/s   39.7mb/s   10.791ms/op [0.171ms - 24.412ms]
11.248: IO Summary:  888 ops 88.791 ops/s 0/81 rd/wr  39.7mb/s 11.227ms/op
```

### Filemicro\_rwritedsync.f

```
finish          3261ops     652ops/s   0.0mb/s    0.000ms/op [0.000ms - 0.001ms]
write-file      3262ops     652ops/s   1.3mb/s    1.521ms/op [0.360ms - 2105.543ms]
29.532: IO Summary:  3262 ops 652.317 ops/s 0/652 rd/wr  1.3mb/s 1.521ms/op
```

### Filemicro\_seqread.f

```
seqread-file    44853ops    4485ops/s 4371.6mb/s  0.221ms/op [0.138ms - 2510.954ms]
33.897: IO Summary: 44853 ops 4484.658 ops/s 4485/0 rd/wr 4371.6mb/s 0.221ms/op
```

#### Filemicro\_seqwrite.f

```
finish          0ops          0ops/s    0.0mb/s    0.000ms/op [0.000ms - 0.000ms]
write-file      444ops         44ops/s   44.3mb/s   22.471ms/op [3.680ms - 24.810ms]
11.300: IO Summary:  444 ops 44.395 ops/s 0/44 rd/wr  44.3mb/s 22.471ms/op
```

#### Filemicro\_statfile.f

```
statfile1      112714ops    11271ops/s  0.0mb/s   1.755ms/op [0.061ms - 10.018ms]
168.727: IO Summary: 112714 ops 11270.545 ops/s 0/0 rd/wr  0.0mb/s 1.755ms/op
```

#### Filemicro\_writesync.f

```
finish          14ops          1ops/s    0.0mb/s    0.001ms/op [0.000ms - 0.006ms]
sync-file       14ops          1ops/s    0.0mb/s   32.960ms/op [23.206ms - 145.686ms]
append-file     15226ops    1522ops/s  11.9mb/s   0.621ms/op [0.179ms - 3.420ms]
11.280: IO Summary: 15240 ops 1523.843 ops/s 0/1522 rd/wr  11.9mb/s 0.651ms/op
```

#### Webserver.f

```
appendlog      1449ops    145ops/s   1.1mb/s   8.365ms/op [0.317ms - 64.755ms]
closefile10    1399ops    140ops/s   0.0mb/s   2.200ms/op [0.078ms - 7.810ms]
readfile10     1399ops    140ops/s   2.2mb/s   3.726ms/op [0.127ms - 13.845ms]
openfile10     1403ops    140ops/s   0.0mb/s  27.577ms/op [2.791ms - 76.815ms]
closefile9     1403ops    140ops/s   0.0mb/s   2.151ms/op [0.073ms - 7.715ms]
readfile9      1403ops    140ops/s   2.1mb/s   3.924ms/op [0.106ms - 14.556ms]
openfile9      1407ops    141ops/s   0.0mb/s  27.865ms/op [4.029ms - 61.863ms]
closefile8     1407ops    141ops/s   0.0mb/s   2.178ms/op [0.071ms - 9.028ms]
readfile8      1407ops    141ops/s   2.2mb/s   3.868ms/op [0.101ms - 13.931ms]
openfile8      1413ops    141ops/s   0.0mb/s  27.304ms/op [3.617ms - 73.206ms]
closefile7     1414ops    141ops/s   0.0mb/s   2.129ms/op [0.077ms - 7.272ms]
readfile7      1414ops    141ops/s   2.1mb/s   3.787ms/op [0.147ms - 16.676ms]
openfile7      1418ops    142ops/s   0.0mb/s  26.949ms/op [3.449ms - 57.598ms]
closefile6     1419ops    142ops/s   0.0mb/s   2.153ms/op [0.083ms - 9.235ms]
readfile6      1419ops    142ops/s   2.2mb/s   3.744ms/op [0.162ms - 13.687ms]
openfile6      1424ops    142ops/s   0.0mb/s  27.247ms/op [3.435ms - 72.481ms]
closefile5     1425ops    142ops/s   0.0mb/s   2.218ms/op [0.066ms - 9.082ms]
readfile5      1425ops    142ops/s   2.2mb/s   3.987ms/op [0.149ms - 13.474ms]
openfile5      1427ops    143ops/s   0.0mb/s  27.231ms/op [3.997ms - 67.959ms]
closefile4     1427ops    143ops/s   0.0mb/s   2.166ms/op [0.077ms - 8.009ms]
readfile4      1427ops    143ops/s   2.2mb/s   3.917ms/op [0.121ms - 13.664ms]
openfile4      1430ops    143ops/s   0.0mb/s  27.162ms/op [3.688ms - 68.930ms]
closefile3     1430ops    143ops/s   0.0mb/s   2.137ms/op [0.077ms - 8.985ms]
readfile3      1430ops    143ops/s   2.2mb/s   4.070ms/op [0.143ms - 13.176ms]
openfile3      1438ops    144ops/s   0.0mb/s  27.058ms/op [3.575ms - 74.976ms]
```

#### Some Takeaways:

Sequential reads/writes perform better than their random counterparts as we expect in a local filesystem as all reads/writes are performed locally.

Workloads with more FS operations like create and mkdir are typically more slower than workloads doing only file operations as filesystem operations in our design always goes to the server.

Also sync operations (for example flushing to disk after every write ) performs poorly because after every write the data has to be flushed to the server as opposed to async operations where it can happen in the background.

## Consistency

The data is flushed to disk after every close operation. We ensure that the client which closes the file last "wins" the write. Here is a small demo showing that :

[https://drive.google.com/file/d/1zqRifms4HjGXSbSg96NgvsrR9uTCtWlm/view?usp=share\\_link](https://drive.google.com/file/d/1zqRifms4HjGXSbSg96NgvsrR9uTCtWlm/view?usp=share_link)

We have tested the functionality using the test case provided and also ran on 2 of our own test cases.

- 1) Close at the end : The 2 clients opens a local copy and do not close the file till the end. We here see that all reads/writes only happen to the local copy and the clients run independently.
- 2) Interleaved close : Both clients start writing to a file simultaneously, one client closes the file before the other (ensured using signals). When the first client opens the file, he sees the contents of the other client and not its own.